

# Architectural Foundations of Modern Football Betting Systems: Ingestion, Persistence, and Predictive Modeling

The construction of a professional-grade football betting system requires a synthesis of high-performance data engineering, relational database optimization, and mathematically rigorous statistical modeling. For a foundational "Batch 1" implementation, the objective is to establish a robust pipeline that transforms raw data from accessible API sources into actionable, calibrated probabilities while maintaining strict backtesting integrity. This report examines the technical requirements for building such a system, focusing on the selection of data providers, the design of resilient extract-transform-load (ETL) architectures, the implementation of sophisticated predictive models like the Dixon-Coles distribution alongside machine learning alternatives, and the enforcement of protocols to eliminate inferential biases during strategy validation.<sup>1</sup>

## Analysis of the Football Data API Ecosystem for Prototyping

The selection of a data provider for the initial phase of a betting system is governed by the tension between data depth, update frequency, and the financial constraints of early-stage development. Free and freemium API tiers in 2025 and 2026 serve as the primary gateway for system designers to validate their modeling hypotheses before scaling to enterprise-grade feeds.<sup>1</sup>

### Comparative Technical Analysis of Free Tier Providers

A comprehensive review of the current market identifies three primary providers suitable for a Batch 1 deployment: API-Football, Football-Data.org, and Sportmonks. These services are characterized by their maturity, documentation quality, and developer-friendly integration paths.<sup>1</sup>

Provider	Access Model	Daily Request Limit	League Coverage (Free)	Key Data Points Available	Update Latency
API-Football	Freemium <sup>6</sup>	100 requests	All competition	Pre-match/in-play	15 seconds <sup>7</sup>

		per day <sup>6</sup>	s (season restricted) <sup>6</sup>	odds, lineups, injuries, predictions <sup>6</sup>	
<b>Football-D ata.org</b>	Free Forever <sup>10</sup>	10 requests per minute <sup>11</sup>	12 top European/ World leagues <sup>11</sup>	Fixtures, results, league tables <sup>11</sup>	Delayed for free tier <sup>11</sup>
<b>Sportmonk s</b>	Free Plan / Trial <sup>14</sup>	180 requests per hour <sup>14</sup>	Danish Superliga, Scottish Premiership <sup>14</sup>	Standard stats, head-to-he ad, team squads <sup>14</sup>	Real-time <sup>14</sup>
<b>TheSports DB</b>	Community -Driven <sup>1</sup>	100 requests per minute <sup>5</sup>	Broad multi-sport coverage <sup>1</sup>	Historical results, team/player basic info <sup>1</sup>	Varies by community update <sup>5</sup>

The structural differences between these offerings dictate the system's geographical and temporal scope. API-Football provides the most expansive access to various data endpoints, including pre-match and in-play odds even on the free plan, which is critical for systems intended to exploit market movements.<sup>6</sup> However, its daily limit is strict, necessitating a highly efficient ingestion strategy that prioritizes high-value updates.<sup>8</sup> Conversely, Football-Data.org focuses on a small subset of "Tier 1" competitions, such as the Premier League and Champions League, making it ideal for systems targeting high-liquidity markets where data consistency is prioritized over league breadth.<sup>11</sup>

## Data Depth and Feature Availability

Beyond simple scores and fixtures, the modeling foundation requires granular statistics to fuel feature engineering. Sportmonks and API-Football provide extensive player-level and team-level statistics, such as possession percentages, shot accuracy, and defensive pressures.<sup>6</sup> For Batch 1, the inclusion of historical odds is non-negotiable, as it allows the system to compare its predicted probabilities against the market's implied probabilities.<sup>3</sup> While free tiers often limit the depth of historical data—frequently restricting users to the current or last season—this is sufficient for establishing the baseline predictive accuracy of a new model.<sup>4</sup>

# Resilient ETL Architecture for Distributed Environments

Data ingestion in the sports betting domain is frequently plagued by network instability, API rate limiting, and the volatile nature of match schedules. A resilient ETL architecture must be stateful, idempotent, and capable of recovering from transient failures without human intervention.<sup>23</sup>

## Implementing Intelligent Retry Logic

Traditional retry loops that use fixed intervals often lead to "thundering herd" problems, where multiple failed ingestion threads all attempt to reconnect simultaneously, potentially triggering firewall blocks or cascading server failures.<sup>24</sup> The industry standard for mitigating this is the use of exponential backoff with jitter, implemented in Python through libraries such as Tenacity.<sup>24</sup>

The mathematical wait time  $W_n$  for attempt  $n$  can be modeled as:

$$W_n = \min(W_{\max}, W_{\text{base}} \cdot 2^n) + \text{Uniform}(0, \text{jitter})$$

This approach ensures that requests are spread out over time, increasing the likelihood of success as the API provider's temporary load subsides.<sup>24</sup> Tenacity allows for the differentiation between transient errors (e.g., HTTP 503 or network timeouts) and permanent errors (e.g., HTTP 404 or 401), ensuring that the pipeline only retries when recovery is theoretically possible.<sup>24</sup>

## Stateful Ingestion and the Checkpointing Mechanism

For Batch 1, the ETL process must avoid full reloads of historical data, which consume excessive API credits and network bandwidth. Instead, an incremental refresh strategy is required.<sup>23</sup> This paradigm relies on "watermarking"—tracking the last record successfully ingested to define the starting point for the next run.<sup>27</sup>

Ingestion Component	Purpose	Technical Implementation
<b>Checkpoint Table</b>	Maintains persistent state across pipeline runs <sup>28</sup>	PostgreSQL table storing job_name, last_timestamp, and status. <sup>28</sup>

<b>Idempotent Loader</b>	Ensures that re-running a job does not create duplicates <sup>23</sup>	SQL INSERT... ON CONFLICT (id) DO UPDATE or UPSERT logic. <sup>27</sup>
<b>CDC (Source Change)</b>	Detects which records have been updated in the API <sup>23</sup>	Querying API with updated_after filter based on the stored watermark. <sup>26</sup>
<b>Audit Log</b>	Tracks data lineage and processing latency <sup>28</sup>	Metadata fields added to every row: ingested_at, source_api_id.

Checkpointing allows a job to resume precisely where it left off in the event of a failure, a feature that is essential for long-running historical backfills.<sup>23</sup> By storing the MAX(updated\_at) for each table, the system can selectively pull only the delta since the previous execution, optimizing resource utilization.<sup>27</sup>

## Data Normalization and Cross-API Synchronization

A significant challenge in sports data engineering is the semantic inconsistency of team and player names across different providers.<sup>32</sup> "Manchester United" may appear as "Man Utd" or "Manchester Utd FC" depending on the source.<sup>32</sup> A resilient ETL must incorporate a normalization layer using fuzzy string matching algorithms.<sup>32</sup>

The use of RapidFuzz or TheFuzz (formerly FuzzyWuzzy) in Python allows for the calculation of the Levenshtein distance between strings.<sup>33</sup> For sports data, the token\_sort\_ratio is particularly effective as it ignores word order, matching "Paris Saint-Germain" with "Saint-Germain Paris" with high confidence.<sup>35</sup> However, because automated matching can produce false positives (e.g., "Manchester City" vs. "Manchester United"), a canonical mapping table in the database is the preferred long-term solution.<sup>32</sup>

## Relational Persistence: PostgreSQL Schema Design for Sports Analytics

The persistence layer must accommodate the highly relational nature of football data while supporting the time-series requirements of market odds.<sup>38</sup> PostgreSQL is the recommended engine due to its support for ACID compliance, advanced indexing, and table partitioning.<sup>31</sup>

### Core Entity-Relationship Modeling

The schema for a Batch 1 betting system is built around the fundamental entities of matches, teams, and odds snapshots. Separation of these domains is critical to ensure that pre-match

predictions are never accidentally informed by post-match results.<sup>41</sup>

## Static Metadata and Fixtures

The teams and leagues tables act as the source of truth for all IDs. The fixtures table is the central junction, storing the schedule and metadata for every match.<sup>38</sup>

SQL

```
CREATE TABLE teams (
    team_id SERIAL PRIMARY KEY,
    external_id INT UNIQUE, -- API provider ID
    full_name VARCHAR(100) NOT NULL,
    short_name VARCHAR(50),
    country VARCHAR(50)
);

CREATE TABLE fixtures (
    fixture_id SERIAL PRIMARY KEY,
    league_id INT REFERENCES leagues(league_id),
    home_team_id INT REFERENCES teams(team_id),
    away_team_id INT REFERENCES teams(team_id),
    kickoff_time TIMESTAMP WITH TIME ZONE NOT NULL,
    season INT,
    round VARCHAR(50)
);
```

## Temporal Market Data

Betting odds are not static; they represent a moving market that reacts to team news, betting volume, and public sentiment.<sup>22</sup> To perform valid backtesting, the system must store snapshots of these odds as they evolve.<sup>41</sup>

SQL

```
CREATE TABLE odds_snapshots (
    snapshot_id BIGSERIAL PRIMARY KEY,
```

```

fixture_id INT REFERENCES fixtures(fixture_id),
bookmaker_id INT REFERENCES bookmakers(bookmaker_id),
market_type VARCHAR(20), -- '1X2', 'AH', 'O/U'
price_h NUMERIC(10, 3),
price_d NUMERIC(10, 3),
price_a NUMERIC(10, 3),
snapshot_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
is_closing_line BOOLEAN DEFAULT FALSE
);

```

By indexing (fixture\_id, snapshot\_at DESC), the system can efficiently retrieve the latest available odds for any given match at any point in history.<sup>38</sup>

## Performance Optimization for High-Volume Queries

As the system scales to handle thousands of leagues and millions of odds snapshots, standard indexing may become insufficient.<sup>8</sup> Partitioning the odds\_snapshots table by snapshot\_at ranges (e.g., monthly partitions) allows the database to skip entire sections of data during backtesting, significantly improving query speed.<sup>40</sup> Furthermore, for feature engineering—which often requires aggregating the "last \$n\$ matches"—the use of Materialized Views or pre-calculated aggregate tables is advised to reduce the computational overhead of complex joins during model training.<sup>44</sup>

## Mathematical Modeling: Dixon-Coles vs. Machine Learning Paradigms

At the heart of any betting system is the predictive engine. For Batch 1, the choice typically lies between the traditional statistical approach of the Dixon-Coles model and modern machine learning (ML) frameworks such as XGBoost or LightGBM.<sup>3</sup>

### The Dixon-Coles Poisson Refinement

The Dixon-Coles model (1997) is a modification of the independent Poisson distribution, which assumes that the goals scored by each team follow a Poisson process.<sup>47</sup> The fundamental limitation of a basic Poisson model is its inability to account for the dependence between teams and the fact that low-scoring results like 0-0 and 1-1 occur more frequently in reality than a simple model would suggest.<sup>21</sup>

The Dixon-Coles model introduces a dependence parameter  $\rho$  and an interaction function  $\tau$  to correct these probabilities:

$$P(X=x, Y=y) = \tau_{\lambda, \mu}(x, y) \cdot \frac{\lambda^x e^{-\lambda}}{x!} \cdot \frac{\mu^y e^{-\mu}}{y!}$$

$$\frac{e^{\lambda\mu} e^{-\lambda\mu}}{y!}$$

Where  $\lambda$  and  $\mu$  are the expected goals for the home and away teams. The  $\tau$  function shifts probability specifically for the (0,0), (1,0), (0,1), and (1,1) scorelines, providing a more accurate representation of draw likelihood.<sup>46</sup>

Furthermore, Dixon-Coles incorporates a time-weighting component  $\phi$ , where the influence of past matches decays according to a parameter  $\xi$ :

$$\text{Weight} = e^{-\xi \cdot \Delta t}$$

This ensures that a team's performance six months ago is less relevant than their performance last week, allowing the model to adapt to shifts in form or managerial changes.<sup>46</sup>

## Machine Learning Alternatives: XGBoost and LightGBM

Modern sports analytics has seen a shift toward gradient-boosted decision trees (GBDTs). These models do not assume a specific underlying distribution (like Poisson) and can ingest hundreds of diverse features, including non-numerical data like weather or referee assignments.<sup>45</sup>

Model	Strengths	Weaknesses
XGBoost	High predictive power; excellent at capturing non-linear feature interactions. <sup>3</sup>	Prone to overfitting on small datasets; requires careful hyperparameter tuning. <sup>47</sup>
LightGBM	Faster training speeds; efficient memory usage for large-scale datasets. <sup>3</sup>	Less effective than XGBoost for small-batch or sparse data.
Random Forest	Robust to outliers; provides reliable feature importance rankings. <sup>32</sup>	Averaging trees can "smooth out" the extreme tail probabilities needed for high-value bets. <sup>56</sup>

Academic literature consistently shows that while ML models can achieve higher classification accuracy (predicting the winner), they often produce uncalibrated probabilities that are less useful for betting than the outputs of a well-tuned Dixon-Coles model.<sup>3</sup>

## Feature Engineering: Leveraging Academic Consensus

The efficacy of a model is primarily determined by its features. Batch 1 should focus on a subset of indicators that have demonstrated persistent predictive value in recent literature.<sup>57</sup>

## Fundamental and Advanced Indicators

Recent studies (2024-2025) emphasize that offensive accuracy and spatial dominance are the strongest predictors of success in elite football.<sup>59</sup>

Feature	Description	Predictive Power
<b>Expected Goals (xG)</b>	Probability of a shot resulting in a goal based on distance and angle. <sup>53</sup>	Higher predictive reliability than actual goals due to reduced noise. <sup>3</sup>
<b>Expected Threat (xT)</b>	Value assigned to carries and passes moving the ball to dangerous zones. <sup>53</sup>	Identifies teams that create danger even without frequent shots. <sup>53</sup>
<b>Rest Differential</b>	Difference in days since the last match for each team. <sup>45</sup>	Significant predictor of second-half fatigue and defensive lapses. <sup>58</sup>
<b>Defensive Pressures</b>	Number of times a team engages an opponent in possession. <sup>51</sup>	Correlates with both shots created (positive) and shots allowed (negative). <sup>59</sup>

The integration of xG is particularly transformative. While a team may win a match 1-0, their xG might be 0.5 vs. their opponent's 2.3. A system that only looks at results would incorrectly conclude the winning team is stronger, whereas an xG-based system would recognize the winner was "lucky" and adjust future probabilities downward.<sup>3</sup>

## Dimensionality Reduction and Selection

Including too many features often leads to "overfitting to the noise," where the model learns patterns that exist only in historical data and do not generalize to the future.<sup>2</sup> Research using English Premier League data suggests that focusing on approximately 20 significant match-related features (distilled from broader datasets) provides the optimal balance between accuracy and model complexity.<sup>61</sup> Techniques like Principal Component Analysis (PCA) can further help identify underlying patterns in player and team performance metrics.<sup>54</sup>

## Probability Calibration: Transforming Scores into

# Wagers

In the context of 1X2 betting (Home, Draw, Away), a raw model output (such as a classification score or a softmax probability) is often distorted. Calibration is the process of mapping these outputs to the true conditional probability of the event.<sup>56</sup>

## Calibration Strategies for 3-Class Outcomes

Scikit-learn's CalibratedClassifierCV provides the standard framework for this task, utilizing cross-validation to prevent the calibrator from overfitting to the training data.<sup>56</sup>

1. **Sigmoid / Platt Scaling:** Fits a logistic regression model to the classifier's scores. It is particularly effective for maximum-margin methods like SVMs and boosted trees, which tend to show sigmoidal distortions in their output.<sup>56</sup>
2. **Isotonic Regression:** A more flexible, non-parametric approach that can correct any monotonic distortion. However, it is prone to overfitting and should only be used when a large amount of validation data is available.<sup>56</sup>
3. **Temperature Scaling:** Designing specifically for multi-class problems, it scales the logits of a model by a single parameter  $T$  before the softmax function. This maintains the ranking of classes while ensuring the final probabilities are better aligned with reality.<sup>56</sup>

A critical requirement in Batch 1 is ensuring that the "Draw" (X) outcome is correctly calibrated. Standard models often under-represent draws, leading to poor returns.<sup>47</sup> By using a 3-class calibration plot, a researcher can visualize whether the predicted 25% draw probability actually results in a draw 25% of the time.<sup>56</sup>

## Backtesting Integrity: Ensuring Real-World Viability

The ultimate failure of many betting systems is not a lack of predictive accuracy, but a failure of backtesting integrity. Strategies that look profitable in simulation often collapse in live markets because they inadvertently used future information.<sup>2</sup>

### Technical Prevention of Look-Ahead Bias

Look-ahead bias occurs when a model uses data (such as late team news or closing odds) that would not have been available at the moment the bet would have been placed.<sup>2</sup> In the architecture of a Batch 1 system, this must be prevented at the database and signal generation layers.<sup>41</sup>

- **Point-in-Time Data:** Every record must be timestamped with `available_at`. When querying historical data for a match on 2025-05-10, the query must explicitly filter: `WHERE available_at <= kickoff_time - INTERVAL '2 hours'`.<sup>41</sup>
- **Signal Lagging:** Signals should be shifted by one day or the relevant time interval to ensure the decision is based solely on prior-day information.<sup>71</sup>

- **Production Parity:** Using the exact same code for both backtesting and live production ensures that the data transformations and feature calculations are identical.<sup>69</sup>

## Validation Frameworks: Out-of-Sample and Walk-Forward

A robust validation protocol goes beyond a simple train-test split.

- **Out-of-Sample Testing:** The final validation of a strategy must occur on data that was never used for model training or hyperparameter tuning.<sup>2</sup>
- **Walk-Forward Analysis:** This involves testing the strategy on consecutive time segments, moving the training window forward by weeks or months.<sup>2</sup> This approach accounts for the changing dynamics of the football market and helps identify when a model's edge has vanished.<sup>2</sup>

By measuring performance using metrics such as the **Rank Probability Score (RPS)**—which penalizes models more heavily for being confident and wrong—and the **Brier Score**, the developer can select the model that offers the most consistent risk-adjusted returns.<sup>46</sup>

## Strategic Conclusions for Batch 1 Development

The establishment of a data and modeling foundation for a football betting system is a multi-disciplinary effort that requires equal focus on engineering stability and statistical rigor. For the initial deployment, the following technical roadmap is recommended:

1. **API Integration:** Prioritize API-Football for its broad endpoint access and in-play odds, but supplement with Football-Data.org for high-fidelity historical results in major European leagues.
2. **ETL Resilience:** Implement a stateful Python pipeline using Tenacity for retry logic and a PostgreSQL checkpoint table to manage incremental loads. Use RapidFuzz to normalize team names across disparate sources.
3. **Schema Design:** Adopt a normalized PostgreSQL structure that separates static fixture data from time-series odds snapshots to ensure backtesting integrity and avoid look-ahead bias.
4. **Modeling:** Implement the Dixon-Coles model as the baseline statistical engine due to its superior handling of draw probabilities and time-decay. Use XGBoost as a comparative ML benchmark, ensuring all outputs are calibrated via Platt scaling or Temperature scaling.
5. **Validation:** Enforce a strict walk-forward backtesting protocol that factors in transaction costs (slippage and overround) and uses only point-in-time data available prior to kickoff.

By following this architecture, the developer creates a system capable of identifying value—situations where the model's calibrated probability is higher than the implied probability offered by the bookmaker. While no model can guarantee profits in the volatile world of sports, a foundation built on these engineering and mathematical principles provides

the necessary edge to compete in professional markets.<sup>21</sup>

This research report provides the architectural blueprint for the **Profitability Engine** and the integrated **Production System**, moving from the calibrated predictive model (Batch 1) to a governed, value-driven betting strategy.

## 1. Value Bet Identification & Market Analysis

The detection of "value" is the mathematical core of the system. It involves identifying discrepancies between the model's objective probabilities and the market's subjective pricing.

- **Implied Probability and Overround:** Bookmaker odds include a "margin" or "overround," ensuring the sum of implied probabilities for a 1X2 market exceeds 100%. The raw implied probability for any outcome is  $P_{\text{implied}} = \frac{1}{\text{Decimal Odds}}$ . To find the "fair" market price, the overround must be removed, typically by simple normalization:  $P_{\text{fair}} = \frac{P_{\text{implied}}}{\sum P_{\text{implied}}}$ .
- **Expected Value (EV) Formula:** A bet is recommended only when it has a positive expected value. The formula for EV per unit stake is:  
$$\text{EV} = (P_{\text{model}} \cdot (\text{Odds} - 1)) - (1 - P_{\text{model}})$$
  
Simplified, this is  $\text{EV} = (P_{\text{model}} \cdot \text{Odds}) - 1$ .
- **Statistically Significant Thresholds:** To account for model variance and the bookmaker's edge, a trigger threshold of  $\text{EV} > 5\%$  or  $10\%$  is recommended. This "margin of safety" ensures that small errors in model calibration do not result in negative real-world returns.
- **Odds Sourcing:** For a Nigerian developer, **API-Football** or **The Odds API** are efficient for real-time closing and opening lines.<sup>1</sup> Accessing multiple bookmakers allows the system to select the "Best Price," which significantly increases long-term ROI by reducing the effective overround.

## 2. Dynamic Bankroll Management & Staking Strategies

Staking determines how much of the bankroll to risk. Poor staking can bankrupt a highly accurate model.

- **Kelly Criterion:** The Kelly Criterion maximizes the long-term growth of the bankroll by sizing bets relative to the edge ( $\text{EV}$ ) and the odds. The optimal fraction  $f^*$  is:  
$$f^* = \frac{p(b+1) - 1}{b}$$

Where  $p$  is the model's win probability and  $b$  is the decimal odds minus 1.

- **Fractional Kelly:** Full Kelly is notoriously aggressive and can lead to massive bankroll swings if the model's probability  $p$  is even slightly overestimated. A **1/4th Kelly** (staking  $0.25 \cdot f^*$ ) is the industry standard for a conservative but profitable growth profile.<sup>3</sup>

### Staking Implementation (Pseudocode):

Python

```
def calculate_stake(bankroll, p_model, odds, kelly_fraction=0.25):
    edge = (p_model * odds) - 1
    if edge <= 0: return 0 # No bet
    f_star = edge / (odds - 1)
    stake = bankroll * f_star * kelly_fraction
    return min(stake, bankroll * 0.05) # Cap at 5% for safety
```

•

## 3. Bet Execution & Order Management

Due to regulatory complexities and the risk of account "gubbing" (limitations) in Nigeria, a semi-automated approach is recommended.

- **Recommendation vs. Executor:** An autonomous executor (using APIs like Betfair<sup>4</sup>) is technically superior but risky for individual accounts. A **Recommendation Engine** that sends a "Bet Slip" via **Telegram** or **FastAPI** dashboard allows for manual approval.
- **Logging Requirements:** Every recommendation must be logged in the PostgreSQL bets table with the following metadata:

Column	Purpose
timestamp	Time of recommendation for point-in-time auditing. <sup>5</sup>
odds_at_time	Market odds used to calculate EV.

model_prob	The probability used for the decision.
calculated_ev	The edge identified.
status	PENDING, PLACED, SETTLED. <sup>7</sup>

#### 4. Advanced Backtesting & Performance Metrics

Backtesting must simulate the "betting brain" under realistic conditions to prevent overfitting.

- **Event-Driven Simulator:** The engine must loop through matches in chronological order, updating the current\_bankroll after each match's settlement before calculating the next match's stake.
- **Key Performance Metrics:**
  - **Sharpe Ratio:** Measures risk-adjusted return:  $S = \frac{E}{\sigma_p}$ .
  - **Profit Factor:** Ratio of gross profits to gross losses; a value  $> 1.5$  indicates a robust strategy.
  - **Max Drawdown (MDD):** The largest peak-to-trough decline. A rule of thumb is a starting bankroll of at least  $\geq 2 \cdot MDD$ .
- **Monte Carlo Stress Testing:** Apply Monte Carlo simulation to the historical bet sequence, shuffling the order of wins/losses 10,000 times. This identifies the "Risk of Ruin"—the probability that the bankroll hits zero due to a localized sequence of losses.

#### 5. Production System Architecture & Orchestration

The system is built as a series of modular Python services orchestrated by a task scheduler like **Prefect** or **Airflow**.<sup>8</sup>

##### Architecture Components:

1. **Ingestion Service:** Daily fetch of fixtures and closing odds using tenacity for resilience against Lagos network instability.
2. **Inference Service:** Loads the calibrated Batch 1 model to generate  $P_{\text{model}}$  for upcoming matches.

3. **Strategy Service:** Calculates \$EV\$ and stake size using the config.yaml parameters.
4. **Alerting Service:** Dispatches recommendations to the user.

Daily Workflow:

Data Fetch -> Incremental ETL Update -> Feature Gen -> Inference -> EV Calc -> Telegram Alert.

## 6. Monitoring & Model Degradation

Models in sports betting suffer from "concept drift" as team tactics or league dynamics change.<sup>9</sup>

- **Real-Time Dashboard:** A Streamlit dashboard should display current P&L, rolling win rate, and "Calibration Drift" (comparing predicted probabilities vs. actual frequencies).
- **Retraining Protocol:** Automated retraining should be triggered if the **Brier Score** (accuracy of probabilities) exceeds a specific threshold or at the end of every 5-match round.<sup>10</sup>

## 7. Synthesis & Final Blueprint

Sample Core Configuration (config.yaml):

YAML

```
strategy:  
    min_ev_threshold: 0.05  
    kelly_fraction: 0.25  
    max_single_bet_pct: 0.05  
    min_odds: 1.50  
    max_odds: 5.0  
leagues:  
    bankroll: 100000 # Initial currency units
```

Prioritized Implementation Roadmap:

1. **MVP (Weeks 1-2):** Build the event-driven backtester and validate the Batch 1 model's profitability on historical data using Fixed Staking.
2. **Logic Engine (Weeks 3-4):** Implement the Kelly staking and the daily automated ETL pipeline (Prefect + FastAPI).
3. **Refinement (Weeks 5-6):** Add Monte Carlo stress testing and the Streamlit

monitoring dashboard.

4. **Scaling (Ongoing):** Integrate real-time odds comparison across 3+ bookmakers to capture maximum value.

I have completed the research for the profitability engine and system architecture. I've integrated the staking logic, value detection formulas, and a modular production design suitable for your environment. Let me know if you would like to move to Batch 3 or refine any specific architectural component.

# **Architectural Blueprint for Production-Grade European Football Betting Systems: Technical Orchestration, Calibration Monitoring, and Resilient Engineering**

The operationalization of a European football betting system represents a significant shift from theoretical model development to industrial-scale data engineering. The transition requires a multifaceted approach focusing on the reliability of the underlying infrastructure, the statistical rigor of performance monitoring, and the structural integrity of the data processing pipelines. A production-ready system must manage high-frequency data ingestion, complex multi-class probabilistic verification, and fault-tolerant execution to ensure that predictive edges are translated into consistent financial performance. This report provides an exhaustive technical blueprint for constructing such a system, prioritizing refinements in orchestration, multi-class scoring mechanisms, and idempotent processing patterns.

## **Technical Refinements in Workflow Orchestration: Airflow vs. Prefect**

The selection of an orchestration framework is a foundational architectural decision that dictates the developer velocity, resource efficiency, and long-term maintenance burden of the betting system. In the contemporary data engineering landscape, the comparison between legacy platforms such as Apache Airflow and modern, API-driven tools like Prefect reveals a divergence in philosophy regarding infrastructure management and runtime flexibility.

### **Infrastructure Complexity and Resource Overhead**

The infrastructure requirements for Apache Airflow are characterized by a centralized, multi-component architecture that can be heavyweight for individual developers or small quantitative teams. To maintain an Airflow 3 environment, one must orchestrate at least four distinct services: a webserver for the user interface, a scheduler responsible for triggering runs and submitting tasks, a DAG processor for file parsing, and a dedicated database such as PostgreSQL for metadata storage. This architecture often requires careful management of `airflow.cfg` files, environment variables, and Docker Compose orchestration, which can extend setup times from hours to days. The official documentation recommends allocating a minimum of 4GB of RAM to the Docker environment, with 8GB suggested for production stability. Furthermore, Airflow lacks

native support for Windows, forcing developers on that platform to utilize virtual machines or the Windows Subsystem for Linux (WSL), which adds further resource consumption and configuration friction.

Prefect, by contrast, adopts a "zero infrastructure" philosophy for local development. It does not require a persistent scheduler, database, or webserver to execute workflows; instead, flows run as standard Python processes. This allows developers to iterate quickly by running code directly using standard Python tools, avoiding the overhead of framework-specific operators and process boundaries that make debugging complex in Airflow. Prefect's hybrid architecture separates the orchestration control plane from the execution environment. This means the code and data remain within the user's infrastructure, while the state management is handled by an external or local API. Teams migrating from Airflow to Prefect frequently report 60% to 70% infrastructure cost savings, as compute resources are only utilized when workflows are actively executing.

Operational Metric	Apache Airflow	Prefect
Minimum RAM Requirements	4GB - 8GB	Standard Python Overhead
Infrastructure Components	4+ Active Services	0 (Native Execution)
Local Setup Complexity	High (Requires Docker/WSL)	Minimal (Pure Python)
Development Speed	Baseline	Reported 3x Faster
Improvement		
Windows Compatibility	Not Native	Full Native Support
Configuration Strategy	airflow.cfg / Environment	Configuration-as-Code

## Runtime Flexibility and Dynamic Workflows

The rigid nature of Airflow's static Directed Acyclic Graphs (DAGs) can be problematic for football betting systems, where match schedules are dynamic and real-time updates are critical. In Airflow, the order and structure of tasks are predefined, and while features like sensors and triggers exist, they often rely on polling mechanisms that introduce latency. Prefect handles dynamic flows natively, allowing for task generation based on runtime data and conditions. This is particularly advantageous for handling match

postponements or rapid changes in market odds, as the pipeline can adapt without requiring complex DAG manipulations.

Developer velocity is further enhanced in Prefect by its Python-first approach. Workflows are defined using simple decorators (`@flow` and `@task`), meaning they can be executed as standard functions within unit test suites or interactive REPL-style debugging sessions. Airflow historically required more operational scaffolding for testing, as tasks are complex objects that the scheduler must parse regularly. For a developer tasked with managing a high-frequency betting system alone, the ability to run tasks individually outside of their flows without spinning up a full service stack represents a significant reduction in operational friction.

## Total Cost of Ownership and Cloud Hosting Strategies

For small to medium-scale betting systems, the financial implications of cloud hosting are a primary constraint. Airflow deployments through managed services such as Astronomer or Amazon Managed Workflows for Apache Airflow (MWAA) can become expensive, with enterprise support often starting around \$45,000 per year. Self-hosting Airflow on virtual machines also entails continuous costs for the 24/7 operation of the scheduler and webserver.

Cloud providers like DigitalOcean and AWS Lightsail offer more predictable pricing models suitable for solo developers. A basic DigitalOcean Droplet with 1GB of RAM and 1 vCPU starts at approximately \$6 per month and includes 1TB of outbound bandwidth, which is a critical consideration for egress-heavy betting pipelines. In contrast, a comparable AWS EC2 t3.micro instance may have a similar base cost but can quickly escalate due to separate charges for storage (EBS) and high data transfer fees (\$0.09 per GB). Prefect's usage-based pricing in its managed cloud service often suits smaller teams better than the tiered support packages of Airflow providers.

Cloud Provider	Instance Type	RAM / CPU	Monthly Cost	Base	Bandwidth Allowance
DigitalOcean	Droplet	1GB / 1 vCPU	\$6.00		1TB Included
AWS Lightsail	VPS	1GB / 1 vCPU	\$5.00		Tiered/Predictable

AWS EC2	t3.micro	1GB / 2 vCPU	~\$8.00 - \$10.00	100GB (Egress) Free
Hetzner Cloud	VM	2GB / 2 vCPU	~\$5.00	20TB Included
Azure	B1s	1GB / 1 vCPU	~\$10.00 - \$12.00	Region Dependent

## Multi-Class Brier Score: A Framework for Probability Verification

In the competitive landscape of European football betting, a model's utility is defined not by its categorical accuracy but by the calibration of its probabilistic forecasts. The Brier Score is a strictly proper scoring rule that assesses the accuracy of these probabilities, penalizing both incorrect predictions and overconfidence. For football markets involving three outcomes—Home Win (1), Draw (X), and Away Win (2)—the multi-class extension of the Brier Score is required to verify model performance.

### Mathematical Definition and Application to 1X2 Markets

The multi-category Brier Score is defined as the mean squared difference between predicted probabilities and actual outcomes across all possible classes. Mathematically, for  $N$  matches and  $R$  outcomes (where  $R=3$  for 1X2 markets), the score is calculated as:

$$BS =$$

$$\frac{1}{N} \sum_{t=1}^N$$

$$t=1$$

$$\sum$$

$$N$$

$c=1$

$\sum$

$R$

$(f$

$tc$

$-o$

$tc$

)

2

In this equation,  $f$

$tc$

represents the forecast probability assigned to outcome  $c$  for match  $t$ , and  $o$

$tc$

is a binary indicator that takes the value 1 if the outcome occurred and 0 otherwise. This formulation ensures that the model is rewarded for assigning high probability to the realized outcome while being penalized for the probability assigned to the other categories.

While binary Brier Scores range from 0 to 1, the multi-class version for 1X2 markets ranges from 0 to 2. A score of 0 reflects perfect accuracy (assigning a 1.0 probability to the actual winner), whereas a score of 2 represents perfect inaccuracy (assigning 1.0 to an outcome that did not occur). For example, if a model predicts a 60% chance of a home win, a 25% chance of a draw, and a 15% chance of an away win, and the home team wins, the calculation for that match is:

(0.60–1.0)

2

+(0.25–0.0)

2

+(0.15–0.0)

2

$$=0.16+0.0625+0.0225=0.245$$

If the match had resulted in an upset away win, the score would increase significantly, reflecting the model's surprise:

(0.60–0.0)

2

+(0.25–0.0)

2

+(0.15–1.0)

2

$$=0.36+0.0625+0.7225=1.145$$

### **Decomposition: Reliability, Resolution, and Uncertainty**

A comprehensive analysis of a betting model requires decomposing the Brier Score into its constituent parts to identify specific areas for improvement. The three-component Murphy decomposition (1973) partitions the score into Uncertainty, Reliability, and Resolution.

1. **Uncertainty (UNC):** This represents the inherent unpredictability of the league's outcomes. It is the variance of the binary outcomes and is independent of the

model's predictions. Uncertainty is maximized when all three outcomes are equally likely (0.33 each).

2. **Reliability (REL) or Calibration:** This measures the degree to which the model's predicted probabilities correspond to real-world frequencies. A model is perfectly reliable if, over many matches where a 70% win probability was forecast, the home team wins exactly 70% of the time. Lower reliability values are better, with 0 indicating perfect calibration.
3. **Resolution (RES):** This term quantifies how much the model's conditional probabilities deviate from the long-term climatological average. A model with high resolution is one that can successfully distinguish between high-probability and low-probability events. Higher resolution values are preferred.

The relationship can be expressed as:  $BS=REL-RES+UNC$ . This insight is vital for a betting system because a model might be well-calibrated (low REL) but have poor resolution (low RES), meaning it provides little "value" beyond the league's base rates.

### Benchmarking and Skill Scores

In practice, a Brier Score must be interpreted relative to a benchmark. A common reference is the "fence-sitter" model, which assigns a uniform probability of 0.33 to each outcome. For a 1X2 market, this reference model yields a Brier Score of  $0.67(2\times(1/3)^2)$

$2\times(1/3)^2$

$+2\times(1/3)\times(1/3)$

$=0.222+0.444$

$=0.666$ . A predictive model must consistently score below this threshold to demonstrate any skill.

The Brier Skill Score (BSS) provides a standardized measure of improvement over a reference model:

$$BSS=1-\frac{BS}{reference}$$

$BS$

*reference*

*BS*

*forecast*

A BSS of 1 indicates perfect forecasts, while a BSS of 0 indicates no improvement over the reference. Negative values suggest the model is performing worse than the benchmark. For professional betting systems, the benchmark is often set as the closing odds of "sharp" bookmakers like Pinnacle, as these incorporate the most efficient market information.

Outcome Scenario	Home Prob	Draw	Away	Actual Result	Brier
	Prob	Prob	Prob		Score
Expected Result	0.8168	0.1195	0.0637	Home Win (1)	0.05
Unexpected Result	0.6984	0.1900	0.1117	Draw (X)	1.16
Random Guess	0.3333	0.3333	0.3333	Any	0.67

## Statistical Monitoring and Practical Threshold-Setting

Machine learning models deployed in betting environments are susceptible to "model decay," a gradual decline in predictive performance caused by shifting league dynamics, squad changes, or tactical innovations. Robust monitoring pipelines must detect these deviations early to trigger automated retraining or system alerts.

### Cumulative Sum (CUSUM) and EWMA Control Charts

Statistical Process Control (SPC) tools such as Cumulative Sum (CUSUM) and Exponentially Weighted Moving Average (EWMA) charts are superior to traditional Shewhart charts for identifying small, persistent shifts in model performance metrics like the Brier Score.

The CUSUM chart aggregates the deviations of the Brier Score from its historical mean. For detecting upward drift (performance degradation), the upper CUSUM statistic is calculated as:

$$S$$

$$hi$$

$$(i) = \max(0, S$$

$$hi$$

$$(i-1) + x$$

$$i$$

$$-\mu$$

$$0$$

$$-K)$$

In this formula,  $x$

$$i$$

is the observed Brier Score,  $\mu$

$$0$$

is the in-control mean, and  $K$  is the "slack" parameter, typically set to  $0.5\sigma$  where  $\sigma$  is the standard deviation. An alarm is triggered when  $S$

$$hi$$

exceeds the decision interval  $H$ , which is often set to  $4\sigma$  or  $5\sigma$ . CUSUM is highly efficient, often reducing initial error alarms by 75% compared to simpler methods.

Alternatively, EWMA charts provide a smoothed view of performance by assigning exponentially decreasing weights to older data. This makes them highly sensitive to recent changes while reducing the impact of noise—a critical feature given the inherent variance in single-match football results. The weight assigned to new information is controlled by the decay factor  $\lambda$ , with common values ranging from 0.1 to 0.4 depending on the desired detection speed for shifts.

Monitoring Parameter	Recommended Value	Strategic Logic
CUSUM Slack ( $K$ )	$0.5\sigma$	Balances sensitivity to small shifts against noise.
CUSUM Threshold ( $(H)$ )	$4\sigma - 5\sigma$	High reliability; minimizes false-positive alerts.
EWMA Smoothing ( $\lambda$ )	0.2	Standard practice; incorporates ~80% historical data.
Target Mean ( $\mu_0$ )	0.48 - 0.52	Based on historical backtesting for top leagues.
In-Control ARL ( $ARL_0$ )	~370	Standard target for industrial-grade monitoring.

## Implementing Automated Response Thresholds

Actionable monitoring requires predefined thresholds that trigger specific system responses. A model degradation of 10% annually is common, but in high-stakes betting, even a slight decline can negate the model's edge over the bookmaker's margin.

Practical thresholds for a betting system include:

- **Warning (Yellow):** Triggered when the EWMA of the Brier Score exceeds the climatological average minus  $1\sigma$ . This alerts the engineering team to monitor for data drift or training-serving skew.
- **Intervention (Orange):** Triggered when the CUSUM statistic exceeds  $3\sigma$ . This should initiate an automated retraining job on the latest available data to attempt recalibration.
- **Critical Stop (Red):** Triggered when the Brier Score exceeds 0.67 or the CUSUM exceeds  $5\sigma$ . Betting execution should be suspended until human intervention can verify the integrity of the input data and the model's validity.

## Resilient Error-Handling and Idempotency Patterns

The volatility of distributed systems and external data providers necessitates a "failure-is-normal" design philosophy. Daily match processing pipelines must be built to handle network failures, API timeouts, and system crashes without corrupting the betting state or duplicating transactions.

### Idempotency and Exactly-Once Semantics

In data engineering, an operation is idempotent if its effect is identical whether it is executed once or multiple times. This is critical for betting systems to avoid duplicate fixture records or double-entry wagers.

Implementation strategies for idempotency include:

1. **Unique Natural Identifiers:** Instead of auto-incrementing surrogate IDs, matches should be identified by stable business keys (e.g., a hash of the date, league, and team names). This ensures that re-inserting the same match results in a collision that can be handled gracefully.
2. **SQL Upserts:** Using `INSERT... ON CONFLICT (id) DO UPDATE` ensures that if a record already exists, it is updated with the latest information rather than creating a duplicate. This is essential for handling odds updates that occur multiple times before a match starts.
3. **Partition Overwrites:** For batch processing of match results, the "delete-write" or "overwrite" strategy is highly reliable. By wiping and replacing a specific day's

data partition, the system ensures no duplicate records remain from a previous partial failure.

## Checkpointing and State Management

Checkpointing allows a pipeline to track its progress and resume from a specific point after a failure, rather than starting from the beginning. This is especially relevant for large backfills or high-volume match ingestion.

A resilient betting pipeline should maintain a metadata table to track the state of match processing:

SQL

```
CREATE TABLE processed_matches (
    match_id VARCHAR(64) PRIMARY KEY,
    process_status  VARCHAR(20), -- 'PENDING', 'SUCCESS',
    'FAILED'
    ingestion_timestamp TIMESTAMP,
    retry_count INT DEFAULT 0,
    error_log TEXT
);
```

By querying this table, the system can skip successfully processed matches and focus only on pending or failed entries, ensuring that "at-least-once" delivery from an API is effectively transformed into "exactly-once" state updates in the database.

## Advanced Error Handling in Orchestrators

Modern orchestrators provide specialized mechanisms for responding to failure. Prefect 3.x utilizes state change hooks to execute code when a task enters a specific state. For example, the `@flow.on_failure` hook can be used to notify a Slack channel with the exact error message and match ID that caused the crash.

Furthermore, Prefect's `allow_failure` annotation allows downstream tasks to continue even if an upstream dependency fails. This is vital for betting systems where the failure to ingest one minor league's odds should not prevent the execution of wagers for more stable markets. The "self-abort" pattern is the idiomatic approach here: a task receives

the failure state of its upstream dependency and programmatically decides to raise `prefect.exceptions.Abort` to skip its own execution without crashing the entire flow.

Resilience Mechanism	Technical Implementation	Purpose
Idempotency Key	<code>Hash(Date + Teams + League)</code>	Prevents duplicate match records.
Upsert Operation	<code>ON CONFLICT DO UPDATE</code>	Updates odds/scores safely on retry.
Checkpointing	Metadata Tracking Table	Resumes processing from last match.
Overwrite Strategy	Partitioned Writes (Daily)	Ensures consistency in batch loads.
State Change Hooks	<code>@task.on_failure</code> / <code>@flow.on_failure</code>	Real-time alerting and incident response.
Conditional Skipping	<code>allow failure + Abort</code>	Decouples league failures in global flows.

## Strategic Synthesis and Future Outlook

The technical blueprint for a European football betting system must integrate these disciplines into a cohesive whole. Orchestration provides the skeleton, ensuring that tasks are executed reliably and with minimal resource waste. The multi-class Brier Score and its accompanying statistical monitoring charts serve as the nervous system, providing constant feedback on the model's calibration and signaling when market conditions have shifted enough to erode the model's edge. Finally, idempotency and resilient error handling act as the immune system, protecting the database's integrity from the inevitable chaos of network partitions and API disruptions.

Looking forward, the integration of ACID-compliant "lakehouse" technologies such as Delta Lake or Apache Hudi can further enhance these systems by providing transactional guarantees previously reserved for traditional databases. The shift towards "non-blocking concurrency control" will allow multiple ingestion processes to

update match data simultaneously without the overhead of complex locking mechanisms, significantly increasing the throughput of the system on high-volume matchdays. Ultimately, the success of a betting system lies in the rigor of its engineering, as the margin for error in sports markets is too thin to tolerate infrastructure instability or silent model degradation. By adhering to the principles of idempotent design and rigorous probabilistic verification, quantitative teams can build platforms that are as resilient as they are predictive.